

Introduction to Garbage Collection

Christian Rinderknecht

Eszterházy Károly Katolikus Egyetem

September 2026

Garbage Collection

The runtimes of many programming languages feature **dynamic memory management**.

The first aspect of that management is **allocation**, that is, the program can create pieces of data (called in this context **objects**) that can outlive the procedure or function that created them, even last the duration of the program itself.

Those objects exist in a memory area called the **heap**. More precisely, a **continuous block** is allocated for each object, whose size depends on the **type** of the object.

Garbage Collection

For example, in Java, we can declare a type for the cells in a linked list of integers as follows:

```
class Cell {  
    int    content;  
    Cell  next;  
}
```

Each cell contains an integer and a **pointer** to another cell. The pointer is implicit in the recursive occurrence of the type name `Cell`. It is either the address of an object in memory, or a special value **nil** (noted `null` in Java). Consequently, each cell requires at least two **words** (one for each field and perhaps more, depending on the data layout).

Garbage Collection

A cell can be allocated with the built-in function `new`, like so:

```
Cell c = new Cell ();
```

A word is the smallest amount of memory that can be addressed and filled by a primitive type or a pointer — 64 bits in modern computers.

In other words, we can imagine memory as an array of words indexed by addresses. Those words may be interpreted as an address or a primitive type.

Then we can modify the contents of the fields it contains, as many times as we wish to. Every call to `new` returns the **address of a new block** drawn from the available memory.

Garbage Collection

Because each allocation removes from available, finite memory, it is necessary for the GC to **free the memory used by useless objects**, which it needs to define and find.


In some programming languages, like C, this task is entirely delegated to the programmers. They have a predefined procedure `free` that frees a memory block:

```
free(c);
```

After it is executed, the block pointed at by `c` is considered available again, so its address can be reused in a future allocation.

Garbage Collection

In C, the premature of an object may lead to errors:



Here, B has been freed. The live object A now contains a **dangling pointer**, and the memory occupied by C has **leaked**: C is not reachable, but it cannot be freed.

Another common mistake is freeing an already freed object (like B above): this is an undefined behaviour and thus a source of **non-determinism**.

Garbage Collection

In Java, there is no equivalent to the `free` procedure of C because the identifying and freeing of useless objects is left to the runtime system.

That part of the runtime is called the **garbage collector** (GC), or 'collector' in short. At regular intervals, it interrupts the program (called hereafter the **mutator**), analyses the heap, find out the useless objects and free them.

This technique was invented by McCarthy, creator of the functional programming language Lisp in 1960. It is widely used in mainstream languages like Java, C#, Haskell or OCaml.

Beyond avoiding the two errors above, a GC may also free many objects at a time, which is more efficient than any sequential, manual deallocation.

Garbage Collection

How does a GC determine what objects are useless, since it has no access to the source code or the programmer (interactively)?

The GC can rely on the following condition:

If, at a certain moment, a program does not have (anymore) a pointer to an object, that object is considered useless.

That condition is not true in C, where pointers can be forged (type casting) from integers, but it holds in Java for example.

Therefore, the memory block of that object can be reused for another, without the program being able to observe that reuse.

Garbage Collection

When, at the current point in the execution of a program, the address of an object is known, that object is said to be **alive**, and **dead** otherwise.

We can define inductively the notion of liveness:

Definition 1

*Both the global and local variables of the program are called **roots**. By definition (base case), the object pointed at by a root is alive.*

Moreover, if an object is alive and one of its fields contains the address of another object, then that object is alive too (inductive case).

What counts as global and local variable depends on the programming language. In Java, the former are the `static` fields (a.k.a. instance fields), and the latter are methods arguments (including `this`) and the variables declared in the method being executed.

Garbage Collection


The heap can now be seen as a **graph**, where the **vertices** correspond to the roots and the objects, and the **edges** model the pointers between objects (that is, if one's block contains a pointer to another's).

Sometimes we will use interchangeably 'block' and 'object' when referring to blocks in the heap. Likewise 'pointer' and 'address'.

The live objects are then the vertices **reachable from the roots**.

Therefore, to determine the live objects at a given time, the GC has to traverse the graph of objects from the roots.

Garbage Collection



Garbage Collection

We started this introduction by referring to 'useful objects', but now we see that the appropriate concept is that of 'live objects', which is the same as 'reachable objects'.

It is obvious that all dead objects are useless, but the reciprocal is not true: a live object might be useless if the program does nothing with it. They are allocated in the heap but cannot be reclaimed. This is a case of **memory leak**. (We saw page 6 another case of memory leak in C.)

Garbage Collection

One of the simplest technique for garbage collection is known as **mark & sweep**.

As the name indicates, it consists of two passes:

1. mark: traversing the *graph* of objects from the roots while marking the visited nodes,
2. sweep: traverse the *whole* heap (the linear piece of memory, not the graph view): all the unmarked objects are freed, and the marks of the others are reset (to mean 'not visited').

Crucially, during the mark & sweep, **the mutator is stopped**.

Garbage Collection

In the following pseudo-code, we will assume that each object o has two fields, $o.left$ and $o.right$, each containing a pointer to another object, or the special value `nil`. Each object also contains a **bit field** denoting the states 'marked' and 'unmarked' (default).

Obviously, the bit marks from the first phase constitute the information carried over to the sweeping phase, but they also play an essential role during marking: they tell the traversal algorithm to not visit a vertex more than once, so the total cost of traversing the graph remains **linear in the number of vertices**.

This is crucial if some objects are shared through distinct paths, in particular if the graph contains cycles.

Garbage Collection

We stated earlier that each object would contain a bit field for the mark. This is the most common solution: put the mark into a **header**, that is, metadata about each object. The downside is that the cost per object becomes (at least) one word, even if only a bit is required and one field of data (that is the worst case: the size of the metadata is the same as that of the data).

There is another choice: a separate table, called a **bitmap**, where the bits are stored in a compact manner: one word will encode many bit marks (64 on modern hardware architectures). This enables getting rid of the object headers and saving a lot of memory, especially if there are many small objects, at the expense of extra time accessing and decoding the table (e.g. mapping object addresses to bit positions).

Garbage Collection

For the marking of objects, any graph traversal can do. For example, let us run a **depth-first search** from each root:

```
DFS(o)
1  if o  $\neq$  nil and o is not marked
2      then mark o;
3      DFS(o.left);
4      DFA(o.right)
```

Note how we use tabulations to suggest blocks of instructions. Finally:

```
MARK()
1  for each root r do DFS(r)
```

Garbage Collection

This program terminates because the bit marks are never set twice nor reset.

Let us count the number of calls to DFS, including the recursive calls. First, there is one call for each root. Second, two recursive calls per object being marked. Therefore, the total number of calls to DFS is $O(r + m)$, where r is the number of roots, and m is the number of marked objects.

Moreover, the flipping of the bit mark is assumed to take a constant time: $O(1)$.

Note that m only counts the live objects: if there are many dead objects, the cost does not change here.

Let us prove something that might seem obvious: namely that, **all reachable objects, and only them, are marked by a call to MARK**.

We proceed by induction on the assertion: 'Object o is reachable from the roots.'

1. In the base case, o is the direct successor of a root, so, by the definition of MARK, one call to $\text{DFS}(o)$ was performed.
2. In the inductive case, o is the direct successor of an object o' that is itself reachable from the roots. The induction hypothesis applies to o' , and says that at least one call to $\text{DFS}(o')$ was made.

Let us consider the first of such calls: o' was unmarked, thus, according to DFS, the test is true, implying a call to $\text{DFS}(o'.\text{left})$, followed (since DFS always terminate on all inputs) by a call to $\text{DFS}(o'.\text{right})$. From our hypotheses, o is either equal to $o'.\text{left}$ or $o'.\text{right}$, therefore a call to $\text{DFS}(o)$ happened.

The reciprocal assertion, that is, if a call $\text{DFS}(o)$ occurs, then o is **nil** or is reachable from the roots, is easily established.

Garbage Collection

The previous pseudo-code uses function calls, but it is better for the collector to not modify the call stack, which is meant for the mutator.

Let us write a single loop instead, one that manages a stack of pointers to be traversed.

MARK()

▷ New definition!

```
1 for each root  $r$  do push  $r$ ;  
2 while the stack is not empty  
3   do pop a pointer  $o$ ;  
4     if  $o \neq \text{nil}$  and  $o$  is unmarked  
5       then mark  $o$ ;  
6         push  $o.\text{left}$ ;  
7         push  $o.\text{right}$ 
```

We could have used a data structure other than a stack, because the order in which the objects are visited does not matter.

We also could avoid pushing **nil** pointers.

If we interpret *o* as a linked list so that *o.left* points to the contents of the cell, and *o.right* points to the next cell, then what is the behaviour of the code above? In particular, what is the size of the stack?

We see that `o.left` is pushed *before* `o.right`, so the next pointer to be popped is `o.right`, and so forth until the end of the list, that is `o.right` is **nil**. So the size of the stack ends up being the same as the list's. If the list is large, this might be a problem.

Instead, we could **push** `o.left` **after** `o.right`. This way, the contents of a cell would be traversed before the next cell, that is, the rest of the list. Therefore the traversal cost (in terms of memory) does not depend on the size of the list.

In a functional programming language, lists are the bread and butter of programmers, therefore a GC traversing the head before the tail of the list is a significant improvement.

MARK() ▷ Second redefinition, more efficient

```
1  for each root  $r$  do push  $r$ ;  
2  while the stack is not empty  
3      do pop a pointer  $o$ ;  
4      if  $o \neq \text{nil}$  and  $o$  is unmarked  
5          then mark  $o$ ;  
6          push  $o.\text{right}$                       ▷ Right first  
7          push  $o.\text{left}$ ;
```

Garbage Collection

As we just saw, the size of the traversal stack can be linear in the size of the reachable data. But how is that stack implemented?

It cannot be allocated in the heap with the mutator's data, as the heap must remain the same during a collection. Instead, it relies on an array whose size is fixed beforehand.

An element of that array then fits a word (interpreted as an address), which is optimal.

Nevertheless, if the heap can contain up to n objects, that array should be of size n , even if not use completely. Moreover, if each block contains at least three words (a header and two pointers), then the length of the stack-array should be a third of that of the heap, which is considerable.

Garbage Collection

One obvious solution is having a **bounded stack** that is not too large. If it becomes full, we drop any new pointer. At the end of the marking, the collector knows that the graph was not traversed in its entirety, because the stack overflow set a flag.

In that case, a linear sweep of the whole heap is performed: whenever a marked object is found to point to an unmarked object, we run the marking algorithm from it. This process is iterated until a fixed point (no new object is marked): all reachable objects have then been marked.

Additionally, in case of a stack overflow, the size of the stack can be doubled for the next collection.

Another solution by Deutsch, Schorr and Waite (DSW), more complex, consists in dispensing with the stack entirely. Instead, the pointers in the heap are modified to keep track of the path of the marking.

Let us illustrate the basic technique behind DSW, called **pointer reversal**, on a singly-linked list containing only primitive types, so each object o has exactly one pointer $o.next$. For instance


$$A \longrightarrow B \longrightarrow C \longrightarrow \mathbf{nil}$$

The collector starts the traversal from A without a stack


The DSW maintains three pointers:

- 1. q points to the current node being visited (to be marked or marked);
- 2. p points to the previous visited node ($p \rightarrow \text{nil}$ if $q \rightarrow A$);
- 3. r points to the rest of the list ($r \rightarrow \text{nil}$ if $q \rightarrow C$).


Initially:




We can mark A and set $A.next$ to point backwards to p :




Then we move p , q and r one step to the right:



We can mark B and set $B.next$ to point backwards to A :




Then we move p , q and r one step to the right:




Garbage Collection

We can mark C and set $C.next$ to point backwards to B :



We are done. To return, we need to reverse the pointers again, this time by swapping p and r :



The general case (trees, graphs) is more complicated, but this gives the gist of it.

Garbage Collection

Let us consider now the second phase: sweeping.


During marking, we considered the graph of objects. Now we look at the heap as an array of words indexed by addresses. (This dual view of the heap, either as a graph or an array, is crucial in understanding garbage collection.)

In such a view of the heap, we see a series of blocks, either of (allocated) objects — in which case they start with a header — or not in use (free blocks).

To understand how to sweep unreachable objects, we first need to know how objects were allocated first.

Garbage Collection

Sequential allocation is a simple method consisting in **partitioning the heap** in two areas: one for the objects (dead or alive) and a **free zone adjacent to it**. The GC has a pointer to the start of the free zone, and another to the upper limit of that zone (to avoid overwriting in case of overflow). A new object of size n is allocated at that address (modulo some alignment padding) as follows:



Garbage Collection

We assume that the collector

- can determine the size of a block by means of `SIZE`,
- and manages a pointer **free** to the start of the free zone, and another, **limit**, to the upper limit of that zone.

```
SEQALLOC(o)                                ▷ o is an object
1  free' ← free + SIZE(o)                  ▷ Read the header of o
2  if free' > limit
3      then SEQALLOC ← nil                  ▷ Insufficient memory
4      else WRITE(o);                       ▷ At the address free
5          free ← free';                     ▷ Bump pointer
6          SEQALLOC ← free
```

This is simple and efficient (locality enables good cache performance).

But, if objects are never moved, deallocations will create free blocks that cannot be reused, because they are not in the free zone: this is called **fragmentation**. We shall come back to this.

Garbage Collection

Another technique does not partition the heap and, instead, assumes that the collector

- can distinguish allocated blocks from those that are free to use;
- can determine the size of each block (allocated and free);
- maintains the set of all free blocks in a **free list**.

Objects already have a one-word header. **Free blocks need a header** too:

1. To distinguish an allocated object from a free block, we store in the first word a **number unique to all free blocks**.
2. In the second word, the **size of the block** is stored.
3. Since free and allocated blocks share the same space (no partitioning of the heap), we need to connect the all free blocks — in a manner reminiscent of how the objects are connected to each other in a graph. So the third word of their header contains a **pointer to the next free block**, making up the free list. The address of the first free block will be held in a global variable.

Garbage Collection

In all, free blocks require three-word header, which is acceptable, as, for instance, we have seen that a simple linked list requires three words of data (the header is metadata).

The free list will require three main operations:

1. a predicate `ISEMPTY()` on whether the free list is empty;
2. a procedure `ALLOCATE(o)` that looks up a free block large enough to contain the object *o*, allocates and initialises the memory, and updates the free list (leaving perhaps a smaller free block) — otherwise fails;
3. a procedure `FREE(o)` that turns an allocated object *o* into a free block that is in turn registered with the free list, perhaps after coalescing it with any contiguous ones.

Garbage Collection

Assuming a free list and the previous primitive operations, we can return to the second phase of the collector: sweeping.

SWEEP()

```
1   $a \leftarrow$  start of the heap  ▷ There is always a block at address  $a$ 
2  while  $a <$  end of the heap
3      do  $n \leftarrow$  SIZE( $a$ )                ▷ Read the header
4          if  $a$  is an allocated object  $o$ 
5              then if  $o$  is marked                ▷  $o$  is alive
6                  then remove the mark ▷ Write the header
7                  else FREE( $o$ )                ▷  $o$  is dead
8       $a \leftarrow a + n$ 
```

Note how we call SIZE(a) *before* calling FREE(o), as the latter could destroy the information needed by SIZE by coalescing the newly freed block with others.

Garbage Collection

We can now write the procedure `NEW` of the GC that is called by the runtime system:

```
NEW(o)
1  if ISEMPTY()
2    then begin
3      MARK();
4      SWEEP();
5      if ISEMPTY()
6        then fail
7        else ALLOCATE(o)
8  else ALLOCATE(o)
```